

SoleMates AI Eval Cheat Sheet

v0.1 — One-page printable. Full 35-term glossary in `eval_cheatsheet.md`.

5 Terms I Use Most

Input / Expected / Actual. What the system sees, what should happen, what did happen. Eval = closing the Expected→Actual gap.

Guardrail. Online metric that triggers immediate action. e.g. *policy_accuracy* runs on every response, blocks/refers on fire.

Improvement flywheel. Offline metric that feeds the next prompt. e.g. *tone* judge on top-20 by score → v3→v3.1 prompt.

Log filtering. Sample by signal, not at random. *score_log* on 100 rows: 1 must / 19 should / 80 low.

Signal-metric divergence. Signals flag a row as interesting, metric says PASS. The 20 top-scored rows all passed policy. Blind spot.

3 Anti-Patterns I Now Recognize

Metric overload. 47-metric dashboard no one reads. Resist until you can name a real failure mode the others miss.

Calibration neglect. Uncalibrated LLM judge > no judge. Test on 50 hand-labeled examples before trusting.

Coverage obsession. 12 hand-picked > 200 generic. Aims for "must not get wrong" + production sample + log-filter long tail.

4-Step Mental Model

Input (user + context) → **Expected** (plain language, debatable) → **Actual** (what really happened) → **Metric** (closes the gap).

If you can't fill in Expected, you don't have a metric. You have a vibe.

If You Only Remember One Thing

Start simple. Add only when you can name the failure mode.

We shipped 4 metrics after 8 chapters. No 5th because no failure mode justified it. The "add a 5th" pressure comes from vendor pitches and 47-metric dashboards.

3 Questions Before Shipping Any Metric

1. How will I know when this metric is lying to me?

Every metric lies sometimes. Code metrics have false +. LLM judges drift. Define the failure mode before shipping.

2. What's the cost per call?

Sub-ms code → online. 5s LLM judge → offline + sampled. The cost difference is 5,000× and the decision is structural.

3. What action does the result trigger?

If no action, why measure? "Tracked weekly" is not an action. "When X crosses Y, ship a new prompt" is.

Things I Still Don't Get (v2 learning list)

- Inter-rater agreement stats (Cohen's kappa, Krippendorff's alpha) — never computed on real data
- Drift detection over time — the v3.1 prompt will drift; no system to alarm on it
- Online guardrail architectures — separate service? pre-commit hook? in-prompt? unclear

Source: `free_courses/ai_evals_for_everyone_labs/chapter10_glossary_of_terms/eval_cheatsheet.md` · v0.1 · PM at SoleMates